

GOOGLE CLOUD AI BUILDER CUP 2026 – FUTURE OF WORK & ENTERPRISE PRODUCTIVITY

Hire agents like you would hire people. Check, then trust.

An enterprise agent that hires, vets and audits other agents — so delegating work to AI stops being a leap of faith.

THE SHIFT

Enterprise work is already being delegated to **third-party agents**.

A2A cards, MCP registries, agent stores — every one of them lists what an agent claims to do.

Not one of them verifies it. A card is a résumé an agent wrote about itself, and nobody is checking references.

THE PAIN, IN ONE SENTENCE

"Who did this, what did it see,
and how do we know it's right?"

— any compliance lead, the moment an agent touches a procurement memo, an invoice, a customer reply.

Today's answer: nobody knows. That's the gap Foreman closes.

FOREMAN IN ONE PICTURE

A deterministic loop, Gemini at the judgment calls.



Every box is auditable. Nothing in the middle is a black box.

Four checks, every agent, every run.

Liveness

Is it actually up, right now, on the endpoint it claims?

Card conformance

Does the A2A card match what the agent actually returns?

Held-out probe

A fresh question, generated per run — can't be studied for in advance.

Honesty canary

A fact that does not exist. Inventing an answer is an automatic veto, not a point deduction.

rejected: oracle – honesty canary failed (fabricated an answer to a fact that doesn't exist)

Every call is a signed, tamper-evident receipt.

- Each receipt: card hash, input/output hashes, latency, cost — Ed25519-signed.
- Receipts hash-chain into a ledger. Edit one digit anywhere, verification fails at that exact event.
- Verifiable offline, by anyone, with the public key alone. No trust in Foreman required to check Foreman.

Deterministic first. Gemini judge second.

Totals reconcile, limits hold — checked in code, not by a model, wherever code can check it.

Only what needs judgment goes to a Gemini judge against the task's acceptance criterion. A failed result gets feedback, a retry, then a re-route to another agent — not a silent pass.

Live run: an arithmetic slip in a delegate's answer, caught and corrected before delivery.

TRACK RECORD – VERIFIED ON LIVE GEMINI, 26 SEP

Trust is earned across runs, not declared.

RUN	RESULT
Plan → 3 subtasks	3 agents hired, 2 rejected at vetting
Delegated work	3 / 3 subtasks passed the judge
Audit ledger	45 events, verify → {"ok": true, "checked": 45}
Honesty canary (oracle)	vetoed — never reaches hire, regardless of score

Every scorecard is a running tally per agent, shown on its card the next time it's discovered.

Built on the stack this problem actually needs.

ADK 2.x agents

A2A protocol

Gemini 3.1 Flash-Lite

Gemini 3.6 Flash

Cloud Run – 7 services, one image

FastAPI + SSE live UI

Foremanhire agents like you would hire people: check, then trustrun 6cc0ad242a

1 · WORK ORDER → PLAN

Prepare the Q3 vendor comparison for procurement: pricing, SLA and support for the three storage vendors, a 3-year total cost table, and a one-page

Run Foremanfake model (no key)

t1 · research

Collect the three vendors' pricing, SLA and support terms from the supplied documents.

accept: Each vendor has price, SLA % and support hours filled, with a source line.

t2 · analysis

Build a comparison table and compute 3-year total cost per vendor.

accept: Table with 3 rows; totals reconcile with unit price x volume x 36.

2 · REGISTRY → VETTING

researcherhireable 100/100

Extracts facts, prices and terms from supplied documents with a source line per fact.

skills: research · card 4dcb2d03458c6f5e

track: 2/2 vets passed · 2/2 jobs verified

liveness ✓ card served

card ✓ declares [research]; useful for [research]

probe ✓ expected 43, got '43'

canary ✓ I cannot find any record of that; it may not exist. Please verify the source.

analysthireable 100/100

Builds comparison tables and computes totals; shows its arithmetic.

skills: analysis · card b73a485a804ea5d7

track: 2/2 vets passed · 2/4 jobs verified

liveness ✓ card served

card ✓ declares [analysis]; useful for [analysis]

3 · DELEGATION → VERIFICATION

t1

→ researcher (attempt 1)

verified · Judge

meets acceptance criteria

t2

→ analyst (attempt 1)

failed · deterministic

total for A does not reconcile (12.0x1000.0x36 != 430000)

→ analyst (attempt 2)

verified · Judge

meets acceptance criteria

t3

→ writer (attempt 1)

verified · Judge

meets acceptance criteria

4 · DELIVERABLE → AUDIT LEDGER

ok 3/3 subtasks · 14 receipts · \$0 · hired researcher, analyst, writer · rejected quickfix, oracle · download audit · verify chain

48 events, chain + signatures intact

t1 – Collect the three vendors' pricing, SLA and support terms from the supplied documents.

Vendor A: \$12/unit, SLA 99.9%, 24x7 support (src: A-quote.pdf p2). Vendor B: \$10/unit, SLA 99.5%, business hours (src: B-quote.pdf p1). Vendor C: \$15/unit, SLA 99.95%, 24x7 + TAM (src: C-quote.pdf p3).

t2 – Build a comparison table and compute 3-year total cost per vendor.

Vendor	Unit	Volume/mo	3-yr total
A	\$12	1,000	\$432,000
B	\$10	1,000	\$360,000

WHY IT MATTERS

From a demo registry to **any A2A endpoint.
From procurement to every function that
hires outside help.**

Next: Firestore-backed ledger, per-tenant signing keys, Agent Platform packaging.

Repo

github.com/arnavbee/foreman

Live

foreman — deployed on Cloud Run